

Introduction to Artificial Intelligence

Adversarial Search


Jianmin Li

Department of Computer Science and Technology
Tsinghua University

Spring, 2024



Outline

- Games
- Perfect Play
- Resource limits and approximate evaluation
- Games of chance




Games vs. search problems

- "Unpredictable" opponent
 - solution is a strategy
 - specifying a move for every possible opponent reply
- Time limits
 - unlikely to find goal, must approximate




History (1)

- The Turk (1770-1854)






W. de Kempelen, del. Ober à Mohel, exécuté à Basile. P. G. Ritz, fo.
Der Schachspieler im Spiele begriffen. Le Joueur d'Échecs tel qu'on le voit pendant le jeu.



History (2)

- Computer considers possible lines of play (Babbage, 1846)
- Algorithm for perfect play
 - Zermelo's Theory, 1912
 - minimax, Von Neumann 1928
 - game of theory: Von Neumann and Morgenstern, 1944



History (3)

- Chess
 - Zuse, 1945; Wiener, 1948; Turing, 1950;
 - Programming a Computer for Playing Chess: Shannon, 1950
 - endgame: D. G. Prinz (1952)
 - full game of standard chess: Bernstein and Roberts, 1958
- Machine learning to improve evaluation accuracy (Checker, Samuel, 1952–57)
- Pruning to allow deeper search
 - alpha-beta: McCarthy, 1956; Newell et al., 1958; Hart and Edwards, 1961; Hart et al. 1972; Knuth and Moore 1975

Types of games

| | deterministic | chance |
|-----------------------|---------------------------------|--|
| perfect information | chess, checkers, go, othello | backgammon monopoly |
| imperfect information | battleships, blind tictactoe | bridge, poker, scrabble nuclear war |

Game tree (2-player, deterministic, turns)

- The initial state
 - How the game is set up at the start
- **PLAYER(s)**
 - Which player has the move in a state
- **ACTIONS(s)**
 - Returns the set of legal moves in a state
- **RESULT(s, a)**
 - The transition model, defining the result of a move

Game tree (2-player, deterministic, turns)

- $\text{TERMINAL-TEST}(s)$
 - TRUE when the game is over and FALSE otherwise
 - terminal states
- $\text{UTILITY}(s, p)$
 - Utility function, also called an objective function or payoff function
 - final numeric value for a game that ends in terminal state s for a player p

Game tree (2-player, deterministic, turns)

MAX (X)


MIN (O)

MAX (X)

MIN (O)

TERMINAL

Utility



Minimax

- Perfect play for deterministic, perfect-information games
- Idea: choose move to position with highest **minimax value**
 - **best** achievable payoff against **best** play

MINIMAX(s) =


$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax

- 2-ply game

$$\text{MINIMAX}(s) =$$

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$



Minimax algorithm

$\text{MINIMAX}(s) =$

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

function MINIMAX-DECISION(*state*) **returns** *an action*

return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$

function MAX-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for each *a* **in** ACTIONS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$

return *v*

function MIN-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

for each *a* **in** ACTIONS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$

return *v*

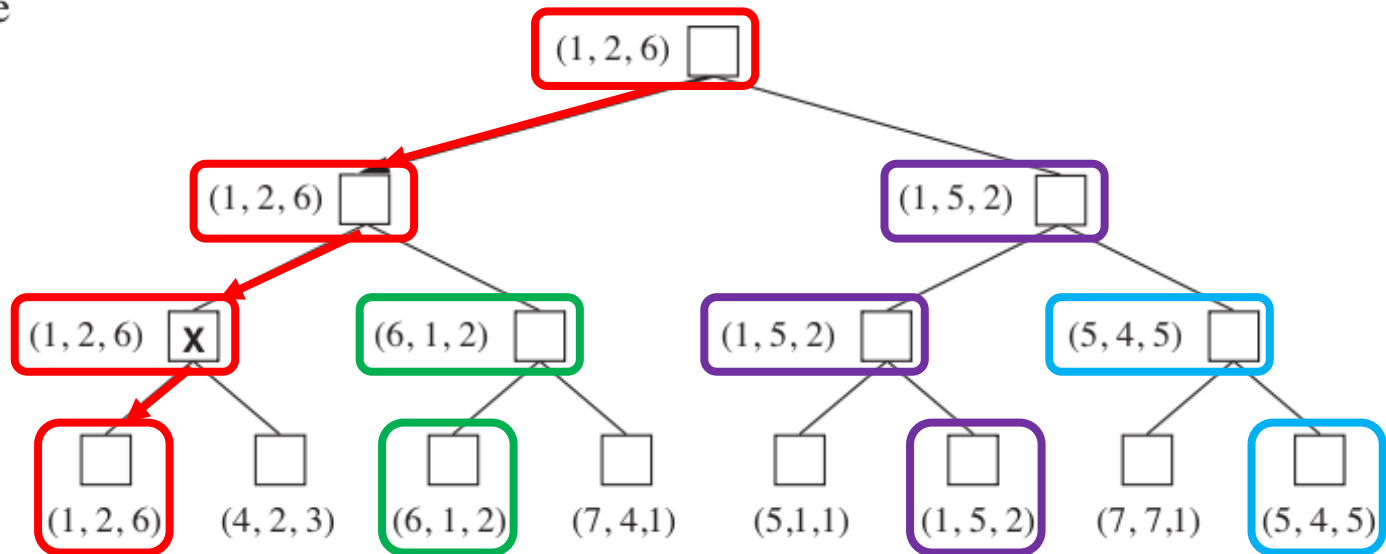
Optimal decisions in multiplayer games

to move
A

B

C

A



MAX for all


Properties of Minimax algorithm

- Complete??
 - Only if tree is finite.
- Optimal??

• But do we need to explore every path?


- $O(b^m)$.
- Space complexity??
 - $O(bm)$
- For chess, $b \approx 35$, $m \approx 100$ for “reasonable” games
 - exact solution completely infeasible

$\alpha - \beta$ pruning example



$$\begin{aligned}
 \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\
 &= \max(3, \min(2, x, y), 2) \\
 &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\
 &= 3.
 \end{aligned}$$

Why is it called $\alpha - \beta$?



- α is the best value (to MAX) found so far off the current path
- If V is worse than α , MAX will avoid it => prune that branch
- β is the best value (to MIN) found so far off the current path

$\alpha - \beta$ algorithm


```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```

Properties of $\alpha - \beta$

- Pruning **does not** affect final result
- **Good** move ordering improves effectiveness of pruning
- With "**perfect ordering**", time complexity = $O(b^{m/2})$
 - doubles solvable depth
- Unfortunately, 35^{50} is still impossible!



Resource limits

- Standard approach

H-MINIMAX(s, d) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN}. \end{cases}$$

- Use **CUTOFF-TEST** instead of **TERMINAL-TEST**
 - e.g., depth limit
 - perhaps add quiescence search
- Use **EVAL** instead of **UTILITY**
 - i.e., evaluation function that estimates desirability of position


Evaluation functions

- For chess, typically linear weighted sum of features


$$\text{Eval}(s) = w_1 f_1 + w_2 f_2 + \dots + w_n f_n$$

- e.g., $w_1 = 9$ with

$$f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$$



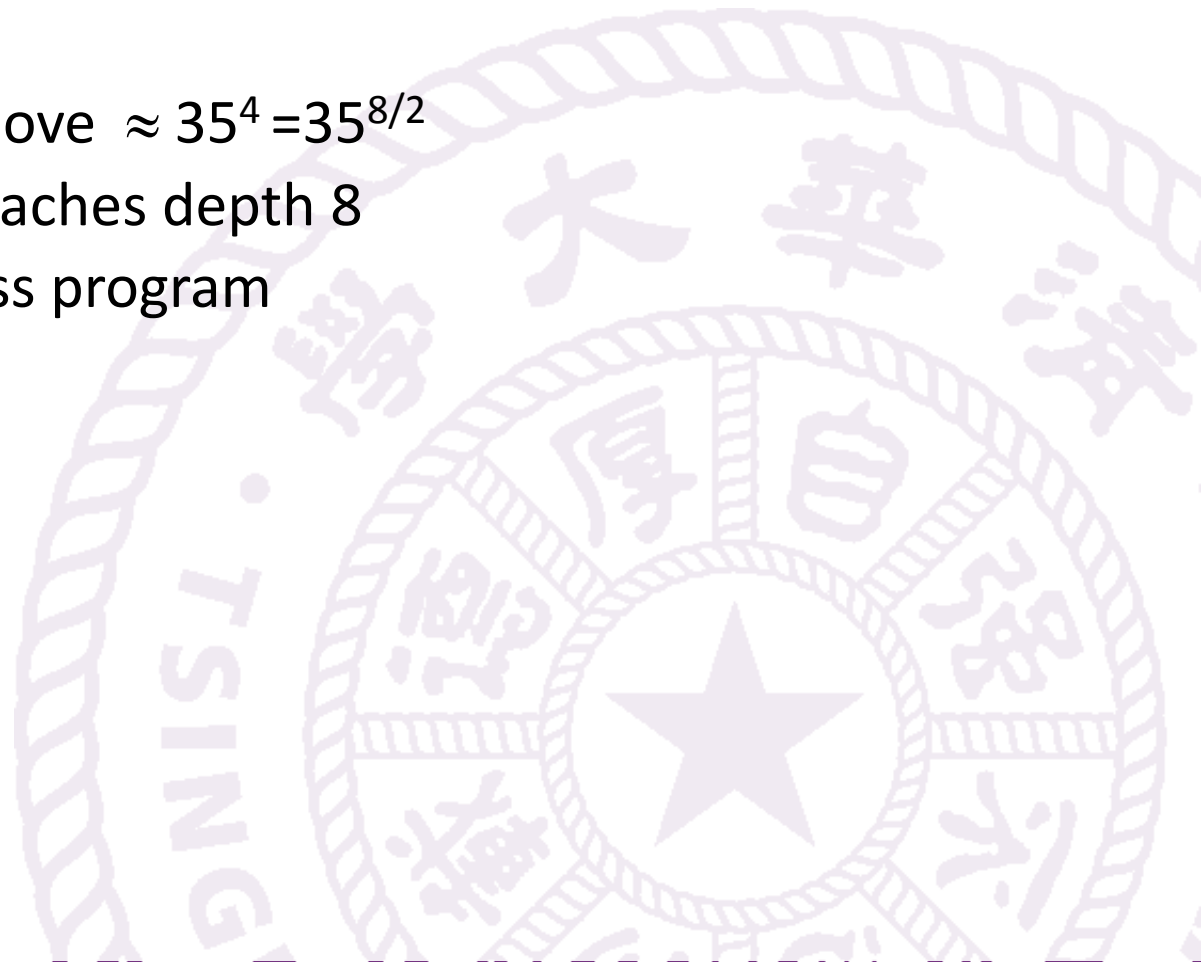
White to move
Black better



White to move
White winning


Resource limits

- Suppose we have 100 seconds, explore 10^4 nodes/second
 - 10^6 nodes per move $\approx 35^4 = 35^{8/2}$
 - $\alpha - \beta$ pruning reaches depth 8
 - pretty good chess program



Digression: Exact values don't matter

- Behavior is preserved under any **monotonic** transformation of EVAL
- Only the order matters:
 - payoff in deterministic games acts as an **ordinal utility** function



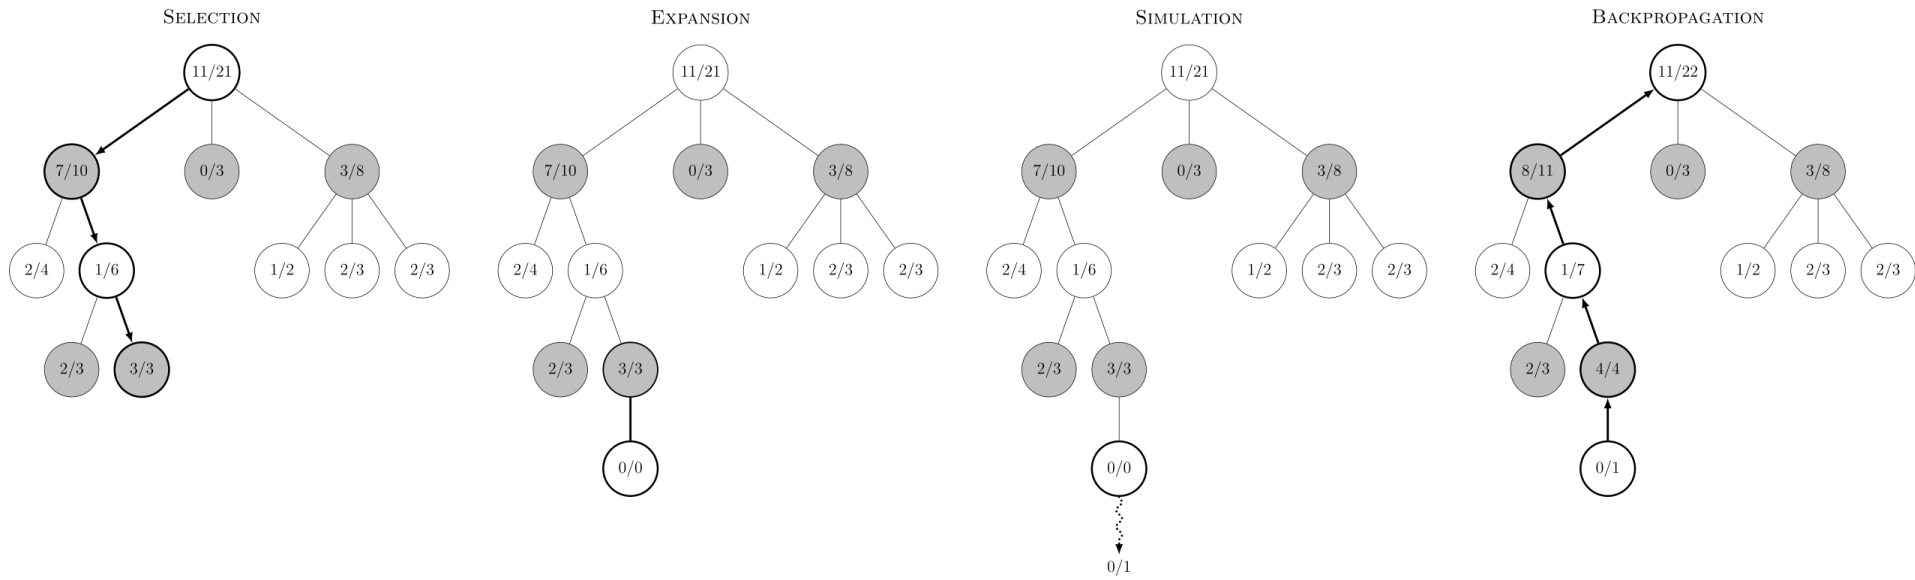
Monte Carlo Tree Search

- Combines the random sampling of traditional Monte Carlo methods with tree searching
- A probabilistic method where random simulations are used to expand the game tree
- Not always find the best move
- Reasonable success at choosing moves that lead to greater chances of winning

Monte Carlo Tree Search

- Selection
- Expansion

- Simulation
- Backpropagation




Monte Carlo Tree Search

- Exploration
 - Promotes exploring unexplored areas of the tree
 - Expand the tree's breadth more than its depth
- Exploitation
 - Stick to one path that has the greatest estimated value
 - Extend the tree's depth more than its breadth

$$UCT(node) = \frac{W(node)}{N(node)} + \sqrt{\frac{\ln(N(parentNode))}{N(node)}}$$

Deterministic games in practice

- Checkers
 - Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used an endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 443,748,401,247 position



Deterministic games in practice

- Chess
 - Deep Blue defeated human world champion Gary Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.



Deterministic games in practice

- Othello
 - human champions refuse to compete against computers, who are too good.




Deterministic games in practice

- Go
 - human champions refused to compete against computers, who were too bad
 - But




Nondeterministic games: backgammon

- $5 \Rightarrow 10, 5 \Rightarrow 11$
- $5 \Rightarrow 11, 19 \Rightarrow 24$
- $5 \Rightarrow 10, 10 \Rightarrow 16$
- $5 \Rightarrow 11, 11 \Rightarrow 16$



Nondeterministic games in general

- In nondeterministic games, chance introduced by dice, card-shuffling
- Simplified example with coin-flipping



Algorithm for nondeterministic games

- EXPECTIMINIMAX gives perfect play
- Just like MINIMAX, except we must also handle chance nodes

EXPECTIMINIMAX(s) =


$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$

MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$


Nondeterministic games in practice

- Dice rolls increase b :
 - 21 possible rolls with 2 dice
- Backgammon ≈ 20 legal moves
 - depth 4 = $20 \times (21 \times 20)^3 \approx 1.2 \times 10^9$
- As depth increases, probability of reaching a given node shrinks
 - value of lookahead is diminished
- TDGAMMON uses depth-2 search + very good EVAL
 - \approx world-champion level



Digression: Exact values DO matter

- Behaviour is preserved only by **positive linear** transformation EVAL
- EVAL should be proportional to the expected payoff



Summary

- Minimax
- H-Minimax
- ExpectMinimax
- Alpha-beta pruning
- Monte Carlo Tree Search



谢谢！

