

Introduction to Artificial Intelligence Beyond Classical Search

Jianmin Li

Department of Computer Science and Technology
Tsinghua University

Spring, 2024

Review

- Assumptions about environment
 - **Observable** => the agent always knows the current state
 - **Discrete** => at any given state there are only finitely many actions to choose from
 - **Known** => the agent knows which states are reached by each action
 - **Deterministic** => each action has exactly one outcome

Review

- The solution to any problem is a fixed sequence of actions
- Tree search algorithm and graph search algorithm
- Search strategy
 - uninformed: BFS, UCS, DFS, IDS
 - informed: Greedy best-first, A*

Today

Relaxing the simplifying assumptions,
getting closer to the **real world**


- Local Search
- Search with Nondeterministic Actions
- Search with Partial Observations

Local Search



Path is irrelevant

- In many optimization problems, path is irrelevant, the goal state itself is the solution
 - n -queens problems
 - IC design
 - communications network optimization
- State space = set of “complete” configurations
 - find optimal configuration, e.g., Travelling Salesperson Problem (TSP)
 - find configuration satisfying constraints, e.g., timetable




Iterative Improvement Algorithms

- Iterative improvement algorithms (local search algorithms)
 - keep a single “current” state
 - try to improve it by moving to neighbors
- Advantages
 - very little memory: usually constant space
 - reasonable solutions in large or infinitely state spaces
 - suitable for online as well as offline search

Example: Travelling Salesperson Problem


- Start with any complete tour, perform pairwise exchanges



- Variants of this approach get within 1% of optimal very quickly with thousands of cities

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts



- Almost always solves n -queens problems almost instantaneously
 - for very large n , e.g., $n = 1\text{million}$

Hill-climbing

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

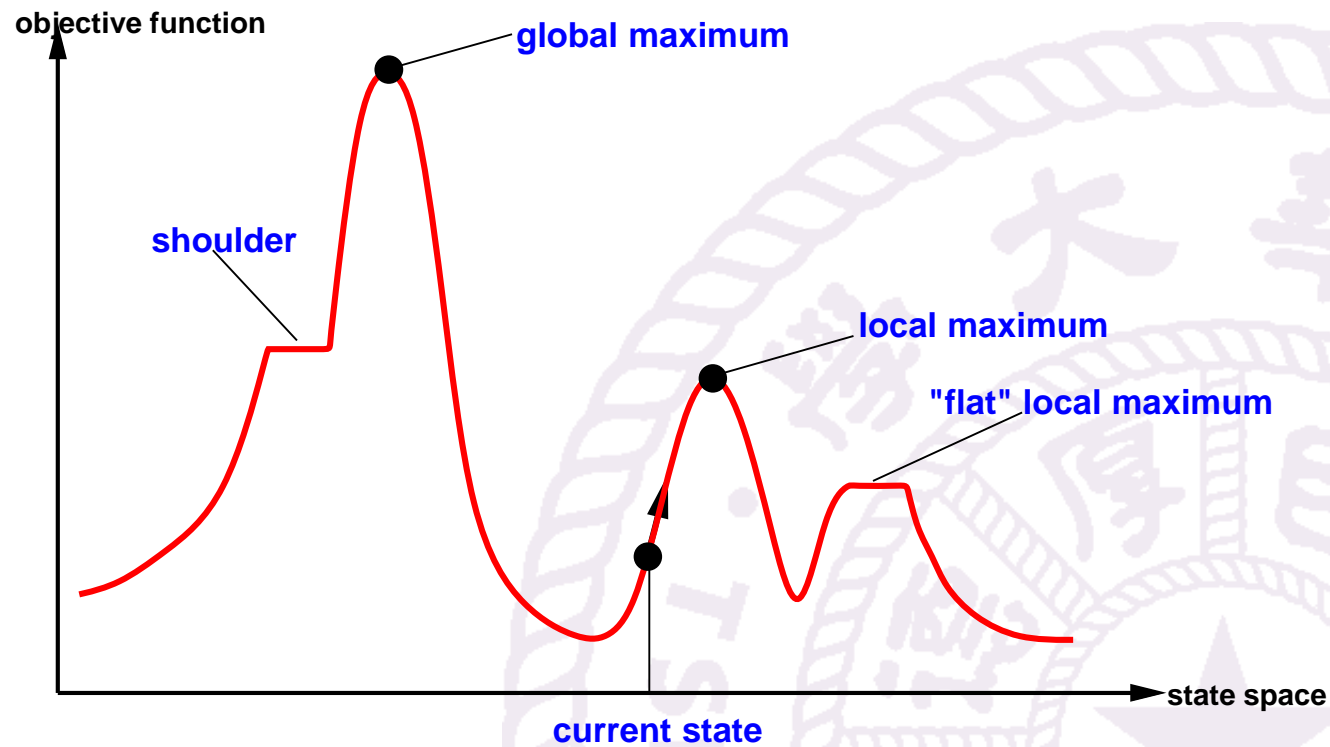
neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE


current \leftarrow *neighbor*

“Like climbing Everest in thick fog with amnesia”

Hill-climbing



Hill-climbing: n -queens




- Complete-state formulation
 - State
 - n queens on the board, one per column
 - Action
 - moving a single queen to another square in the same column
 - each state has $8 \times 7 = 56$ successors.
- Heuristic cost function h
 - the number of pairs of queens that are attacking each other, either directly or indirectly

Question 1

- An 8-queens state. heuristic cost estimate $h = ?$

- A. 0
- B. 1
- C. 2
- D. 3






Hill-climbing: n -queens

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

- An 8-queens state
- heuristic cost estimate
 $h = 17$
- Best successor, $h = 12$


Getting Stuck

- Local maxima
- Ridges
- Plateaux



Improvements

- Random sideways moves
 - escape from shoulders
 - loop on flat maxima
- Random-restart hill climbing
 - overcomes local maxima - trivially complete



Simulated Annealing

- Idea: escape local maxima by allowing some "bad" moves
- gradually decrease their size & frequency

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow$ *schedule*(t)

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ *next*.VALUE – *current*.VALUE

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Properties of Simulated Annealing

- If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Devised by Metropolis et al. (1953) for physical process modeling
- Widely used in VLSI layout, airline scheduling, etc

Local Beam Search

- Keep track of (top) k states rather than just one
 - Start with k randomly generated states
 - At each iteration, all the successors of all k states are generated
 - If any one is a goal state, stop; else select the k best successors from the complete list and repeat.
- Question: k searches run in parallel?

Local Beam Search

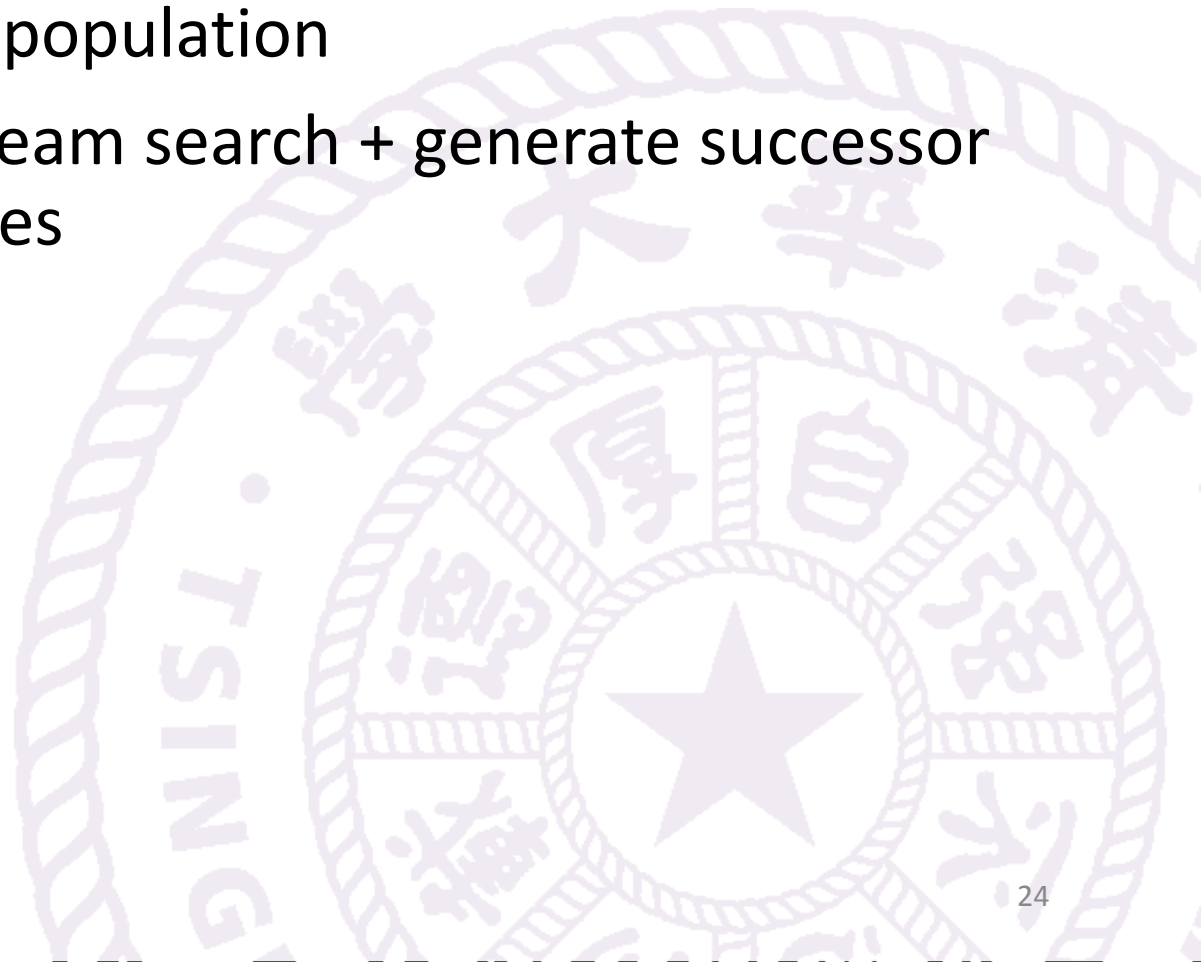
- Problem: all k states end up on same local hill
- Solution idea: choose k successors randomly, biased towards good ones (**Stochastic Beam Search**).
- Close analogy to natural selection

Genetic Algorithms


- **Individual** (i.e. state): each is represented as a string over a finite alphabet (often a string of 0s and 1s)
- **Population**: Start with k randomly generated individuals
- **Fitness function**: evaluation of the “goodness” of a given state
- Produce the next generation of states by **selection**, **crossover**, and **mutation**

Genetic Algorithms

- A successor is generated by combining two parents from the current population
- stochastic local beam search + generate successor from pairs of states



Genetic Algorithms: n -queens



- Fitness function: number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)

$$24/(24+23+20+11) = 31\%$$


$$23/(24+23+20+11) = 29\%$$

$$20/(24+23+20+11) = 26\%$$

$$11/(24+23+20+11) = 14\%$$

Encoding

- GAs require states encoded as strings
- Crossover helps iff substrings are meaningful components



A Genetic Algorithm

function GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

inputs: *population*, a set of individuals

FITNESS-FN, a function that measures the fitness of an individual

repeat

new_population \leftarrow empty set

for $i = 1$ **to** SIZE(*population*) **do**

x \leftarrow RANDOM-SELECTION(*population*, FITNESS-FN)

y \leftarrow RANDOM-SELECTION(*population*, FITNESS-FN)

child \leftarrow REPRODUCE(*x*, *y*)

if (small random probability) **then** *child* \leftarrow MUTATE(*child*)

add *child* to *new_population*

population \leftarrow *new_population*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to FITNESS-FN

function REPRODUCE(*x*, *y*) **returns** an individual

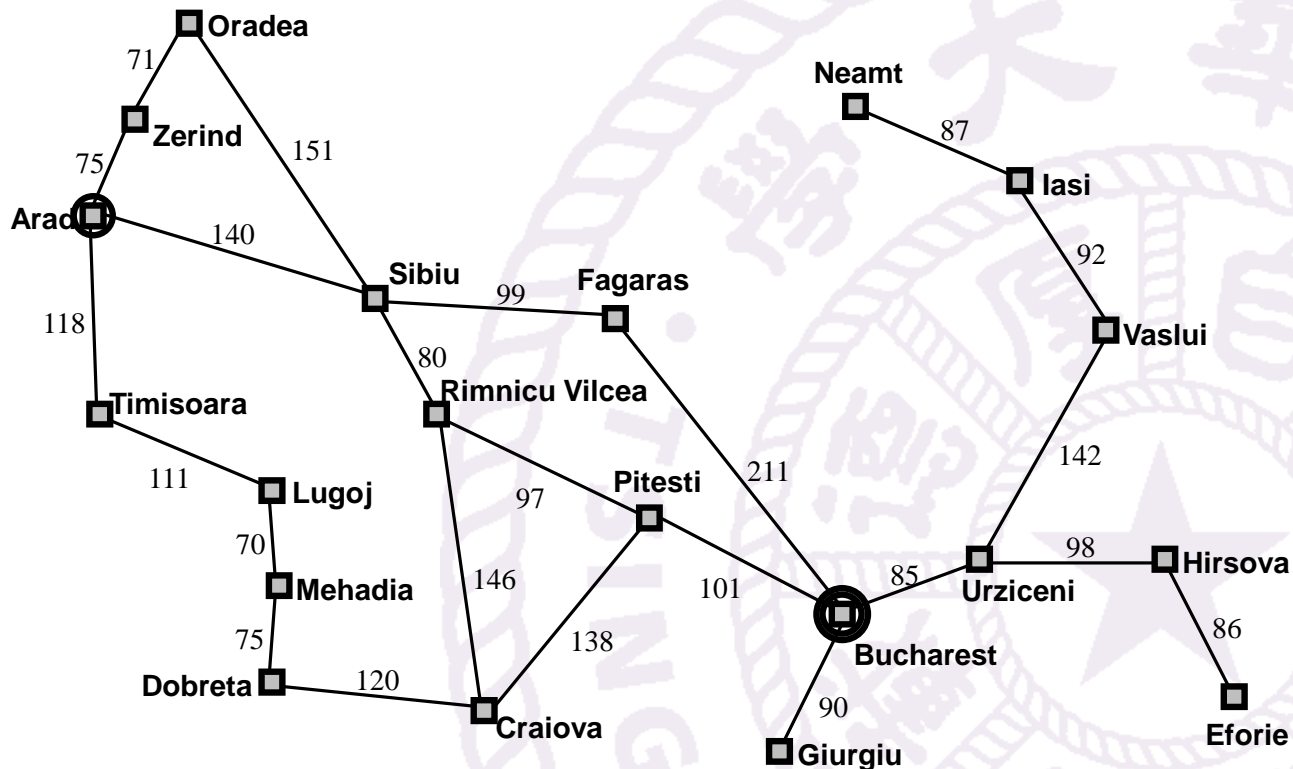
inputs: *x*, *y*, parent individuals

$n \leftarrow$ LENGTH(*x*); $c \leftarrow$ random number from 1 to n

return APPEND(SUBSTRING(*x*, 1, c), SUBSTRING(*y*, $c + 1$, n))

Continuous Search Space

- Three new airports in Romania, such that the sum of squared distances from each city on the map to its nearest airport is minimized



Continuous Search Space

- 6-D state space

$$(x_1, y_1, x_2, y_2, x_3, y_3)$$

- Objective function

- sum of squared distances from each city on the map to its nearest airport

$$f(x_1, y_1, x_2, y_2, x_3, y_3) = \sum_{i=1}^3 \sum_{c \in C_i} (x_i - x_c)^2 + (y_i - y_c)^2$$

- C_i is the set of cities whose closest airport (in the current state) is airport i

Continuous Search Space

- Most real-world environments are continuous
 - a continuous state space

$$S = \{ (x_1, x_2 \cdots, x_N) \mid x_i \in R \}$$

- a continuous object function


$$f(x_1, x_2 \cdots, x_N)$$

- Successor function would return **infinitely** many states!
- Discretization
 - Turns continuous space into discrete space

Gradient

- The gradient of the objective function is a vector
- The gradient gives the magnitude and direction of the steepest slope at a point
- empirical gradient considers $\pm\delta$ change in each coordinate

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_N} \right)$$



the gradient of the function $f(x, y) = (\cos^2 x + \cos^2 y)^2$

Steepest-ascent Hill-climbing

$$x_{k+1} = x_k + \alpha \nabla f$$

- a is step size
 - if a is too small, too many steps are needed
 - if a is too large, the search could overshoot the maximum
- Line search
 - Extending the current gradient direction - usually by repeatedly
 - doubling a - until f starts to decrease again

Newton-Raphson Method

- To find a maximum or minimum of f , we need to find x such that the gradient is zero, i.e. $\nabla f(x) = 0$
- Newton-Raphson method
 - a general technique for finding roots of function, i.e. solving equations of the form $g(x) = 0$
 - computing a new estimate for the root x according to Newton's formula

$$x_{k+1} = x_k - \frac{g(x_k)}{g'(x_k)}$$

- $g(x) = \nabla f(x)$

$$x_{k+1} = x_k - H_f^{-1}(x_k) \nabla f(x_k) \quad H_{\{ij\}} = \frac{\partial^2 f}{\partial x_i \partial x_j}$$

Search with Nondeterministic Actions



Percepts

- In previous slides we assume that
 - the environment is **fully observable** and **deterministic**
 - the agent knows what the effects of each action are
- Therefore, the agent
 - can calculate exactly which state results from any sequence of actions
 - and always knows which state it is in
- Its percepts provide no new information after each action

Percepts


- Percepts become useful in **partially observable** or **nondeterministic** environment
 - In a partially observable environment, every percept helps narrow down the set of possible states the agent might be in, thus making it easier for the agent to achieve its goals
 - When the environment is nondeterministic, percepts tell the agent which of the possible outcomes of its actions has actually occurred

Percepts

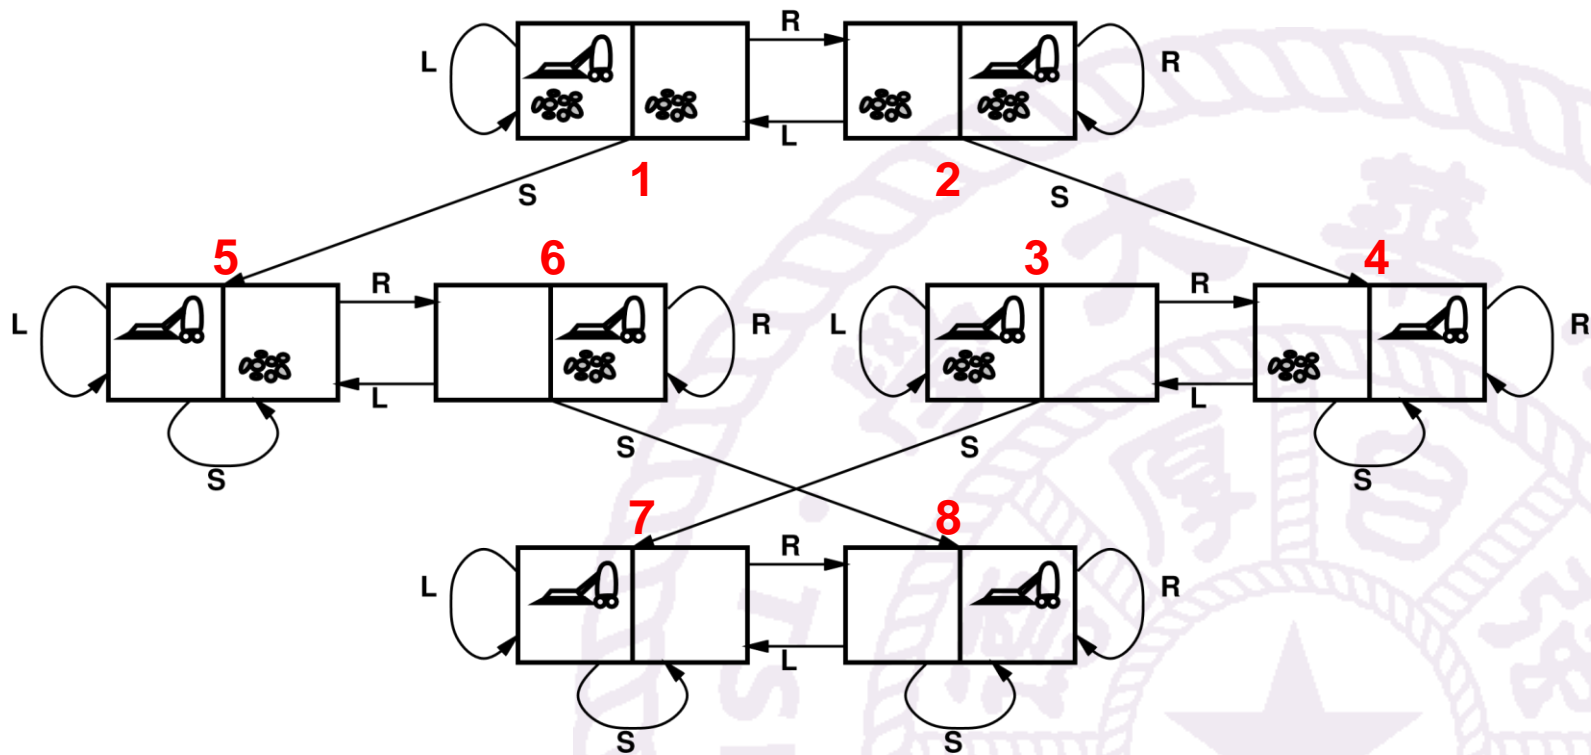
- The future percepts cannot be determined in advance and the agent's future actions will depend on those future percepts
- Solution to a problem is a **contingency plan (or strategy)** rather than a sequence of actions
 - specifies what to do depending on what percepts are received

Deterministic Vacuum World

- actions = {left, right, suck}
- environment is observable, deterministic, and completely known




Deterministic Vacuum World



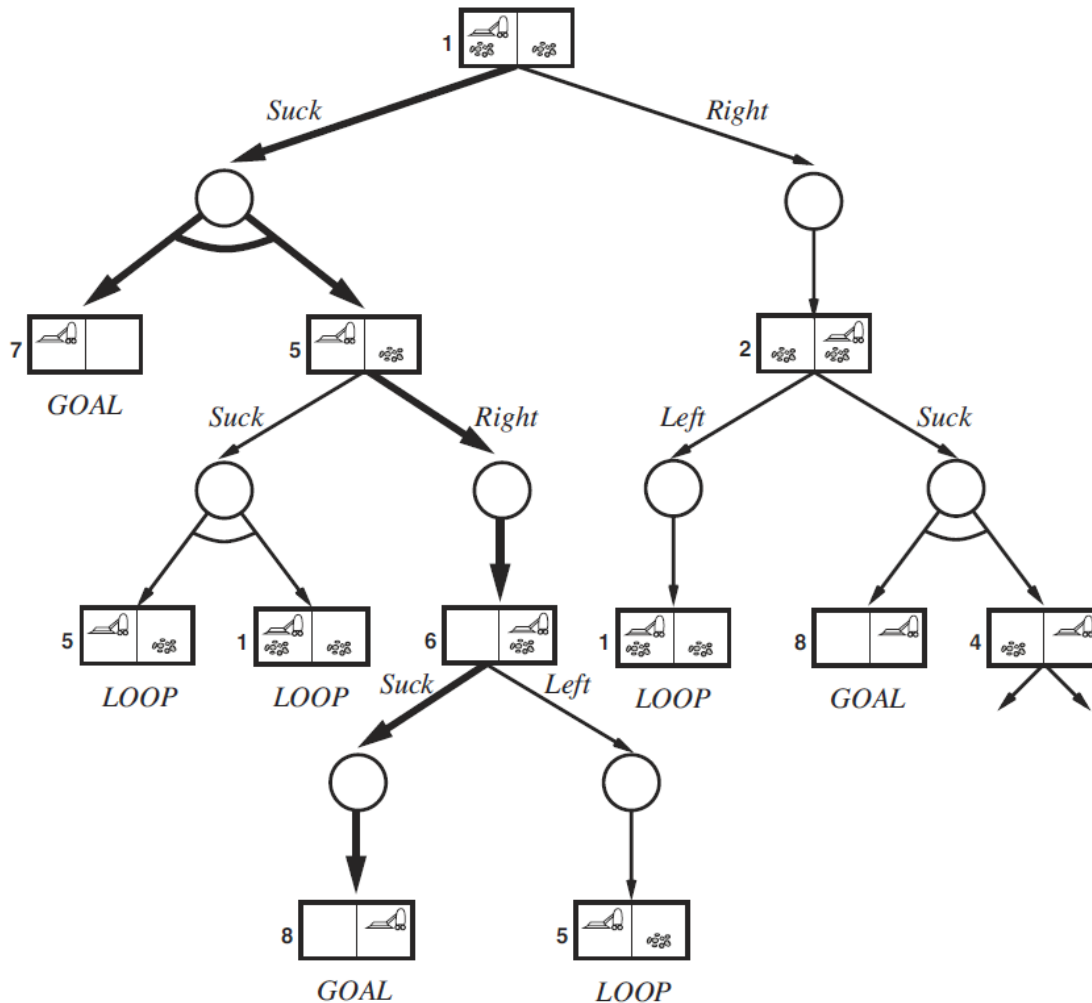
[Suck, Right, Suck]

Erratic Vacuum World

- actions = {left, right, suck}
- When sucking a dirty square, it cleans it and sometimes cleans up dirt in an adjacent square
- When sucking a clean square, it sometimes deposits dirt on the carpet



AND-OR Search Tree



OR nodes

- search possible actions

AND nodes

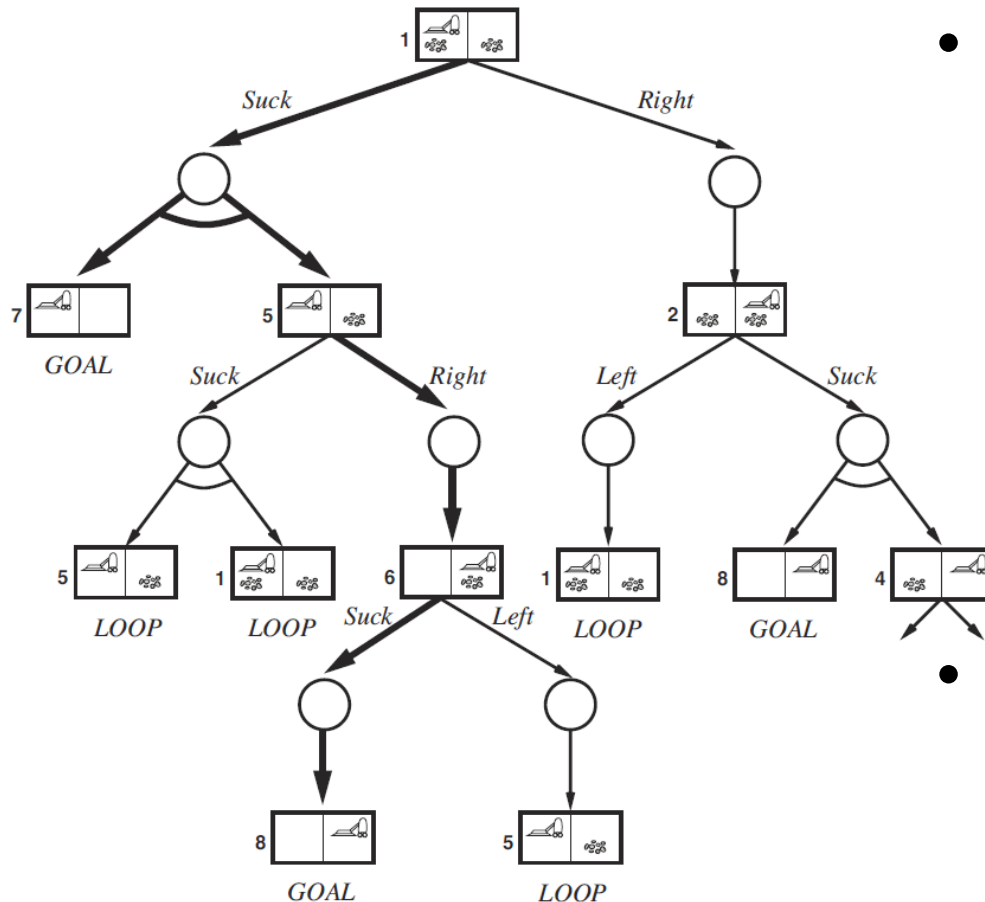
- consider all outcome states


Loops

- terminate search (on that path)

Suck action in state 1 leads to a state in the set {5, 7}, so the agent would need to find a plan for state 5 **and** for state 7.

AND-OR Search Tree



- Solution
 - a goal node at every non-loop leaf
 - specify one action at each of its OR nodes
 - includes every outcome branch at each of its AND nodes
- 
- [Suck, if State = 5 then [Right, Suck] else []]

A Recursive, Depth-first Solution

function AND-OR-GRAPH-SEARCH(*problem*) **returns** *a conditional plan, or failure*
OR-SEARCH(*problem*.INITIAL-STATE, *problem*, [])

function OR-SEARCH(*state*, *problem*, *path*) **returns** *a conditional plan, or failure*
if *problem*.GOAL-TEST(*state*) **then return** the empty plan
if *state* is on *path* **then return failure**
for each *action* **in** *problem*.ACTIONS(*state*) **do**
 plan \leftarrow AND-SEARCH(RESULTS(*state*, *action*), *problem*, [*state* | *path*])
 if *plan* \neq failure **then return** [*action* | *plan*]
return failure

function AND-SEARCH(*states*, *problem*, *path*) **returns** *a conditional plan, or failure*
for each s_i **in** *states* **do**
 *plan*_{*i*} \leftarrow OR-SEARCH(s_i , *problem*, *path*)
 if *plan*_{*i*} = failure **then return failure**
return [**if** s_1 **then** *plan*₁ **else if** s_2 **then** *plan*₂ **else** ... **if** s_{n-1} **then** *plan* _{$n-1$} **else** *plan* _{n}]

- Save path to avoid loops
- Search over **all** the possibilities for an uncertain action outcome
- Solution is contingency plan, dealing with each possible outcome

Search with Partial Observations


Uncertainty Strikes Twice

- Uncertainty in action
 - as above
- Uncertainty in sensing
 - no observation
 - partial observations




Belief State

- Belief state
 - representing the agent's current belief about the possible physical states it might be in



physical states for vacuum



belief states for sensorless vacuum

- To plan a sequence of actions, the agent searches a space of **belief states**, instead of a space of states

Searching with No Observation

- Sensorless problem
 - agent's percepts provide **no information** at all
- Is it possible for a sensorless agent to solve a problem if has no idea what state it's in?
 - Sensorless problems are **quite often solvable**

Searching with No Observation

- Sensorless agent can be surprisingly useful
- they don't rely on sensors working properly
 - many ingenious methods in manufacturing systems
- the cost of sensing is too high
 - doctors often prescribe a broadspectrum antibiotic

Search in Belief State Space

- Search in space of belief states instead of physical states
- In belief-state space, the problem is fully observable
 - the agent always knows its own belief state
- Furthermore, the solution (if any) is always a sequence of actions
 - the percepts received after each action are completely predictable - **they're always empty!**
 - there are no contingencies to plan for
 - This is true even if the environment is **nondeterministic**

Sensorless Problem Definition

- Suppose the underlying physical problem P is defined by
 - $ACTIONS_P$, $RESULT_P$, $GOAL-TEST_P$ and $STEP-COST_P$
- How to define corresponding sensorless problem?

Sensorless Problem Definition

- Belief states
 - The entire belief-state space contains every possible set of physical states
 - **Bad news**
 - If P has N states, then the sensorless problem N has up to 2^N states
 - **Good news**
 - Many may be unreachable from the initial state
- Initial state
 - Typically the set of all states in P
 - In some cases the agent will have more knowledge

Sensorless Problem Definition

- Actions

- The agent is in belief state $b = \{s_1, s_2\}$, but
$$ACTIONS_p(s_1) \neq ACTIONS_p(s_2)$$

- The agent is unsure of what actions are legal
 - If illegal actions have no effect on the environment

$$ACTIONS(b) = \bigcup_{s \in b} ACTIONS_p(s)$$

- If an illegal action might be the end of the world

$$ACTIONS(b) = \bigcap_{s \in b} ACTIONS_p(s)$$

Sensorless Problem Definition

- Transition model

- For **deterministic** actions, the set of states that might be reached is

$$\begin{aligned} b' &= RESULT(b, a) \\ &= \{s' : s' = RESULT_P(s, a), \text{ and } s \in b\} \end{aligned}$$

- For **nondeterministic** actions, the set of states that might be reached is

$$\begin{aligned} b' &= RESULT(b, a) \\ &= \{s' : s' \in RESULTS_P(s, a), \text{ and } s \in b\} \\ &= \bigcup_{s \in b} RESULTS_P(s, a) \end{aligned}$$

Sensorless Problem Definition

- Transition model
 - Prediction step
 - generating the new prediction belief state after the action
$$b' = \text{PREDICT}(b, a)$$
 - With deterministic actions, b' is never **larger** than b
 - With nondeterministic actions, b' may be **larger** than b



Sensorless Problem Definition

- Goal test
 - a belief state satisfies the goal only if **all** the physical states in it satisfy $GOAL-TEST_p$
 - The agent may accidentally achieve the goal earlier, but it won't know that it has done so


Sensorless Problem Definition

- Path cost
 - the cost of taking an action in a given belief state could be one of several values
 - A simple case
 - the cost of an action is the **same** in all states
 - can be transferred directly from the underlying physical problem

Deterministic, Sensorless Vacuum World

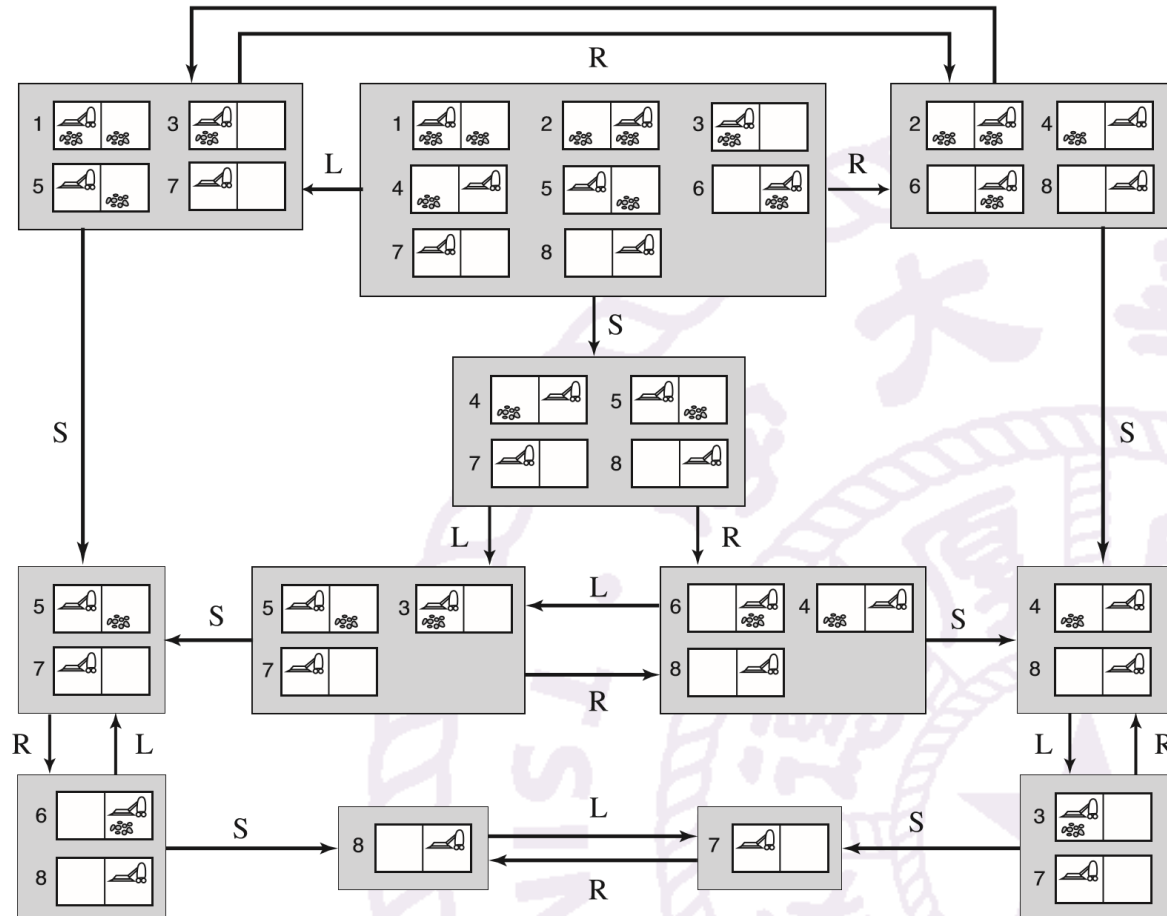
- Assume that the agent knows the geography of its world, but doesn't know its location and the distribution of the dirt.
- Its initial state could be any element of the set of all states

Deterministic, Sensorless Vacuum World



How to reach the goal state (7)?
[right, suck, left, suck]

Sensorless Problem Solving



Apply **any** of the search algorithm we have known

Searching with Partial Observations

- How the environment generates percepts for the agent
 - Local-sensing vacuum world: a position sensor and a local dirt sensor, no global dirt sensor
- If sensing is **deterministic**, a $PERCEPT(s)$ function that returns the percept received in a given state
 - **Fully observable** problem: $PERCEPT(s)=s$ for every state s
 - **Sensorless** problem: $PERCEPT(s)=\text{null}$
- If sensing is **nondeterministic**, then function $PERCEPTS$ that returns a set of possible percepts

Partially Observable Problem Definition

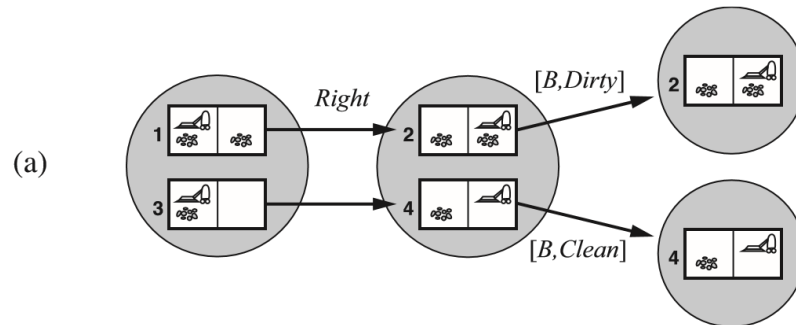
- *ACTIONS*, *STEP-COST*, and *GOAL-TEST* are the same as for sensorless problems
- The transition model is a bit more complicated
 - The **prediction** stage
$$\hat{b} = \text{PREDICT}(b, a)$$
 - The **observation prediction** stage
$$\text{POSSIBLE-PERCEPTS}(\hat{b}) = \{o: o = \text{PERCEPT}(s) \text{ and } s \in \hat{b}\}$$
 - The **update** stage
$$b_o = \text{UPDATE}(\hat{b}, o) = \{s: o = \text{PERCEPT}(s) \text{ and } s \in \hat{b}\}$$

Partially Observable Problem Definition


- Possible belief states resulting from a given action and the subsequent possible percepts:

$$\begin{aligned} & RESULTS(b, a) \\ &= \{b_o : b_o = UPDATE(PREDICT(b, a), o) \text{ and } o \end{aligned}$$

Partially Observable Problem Definition



- a. Deterministic world

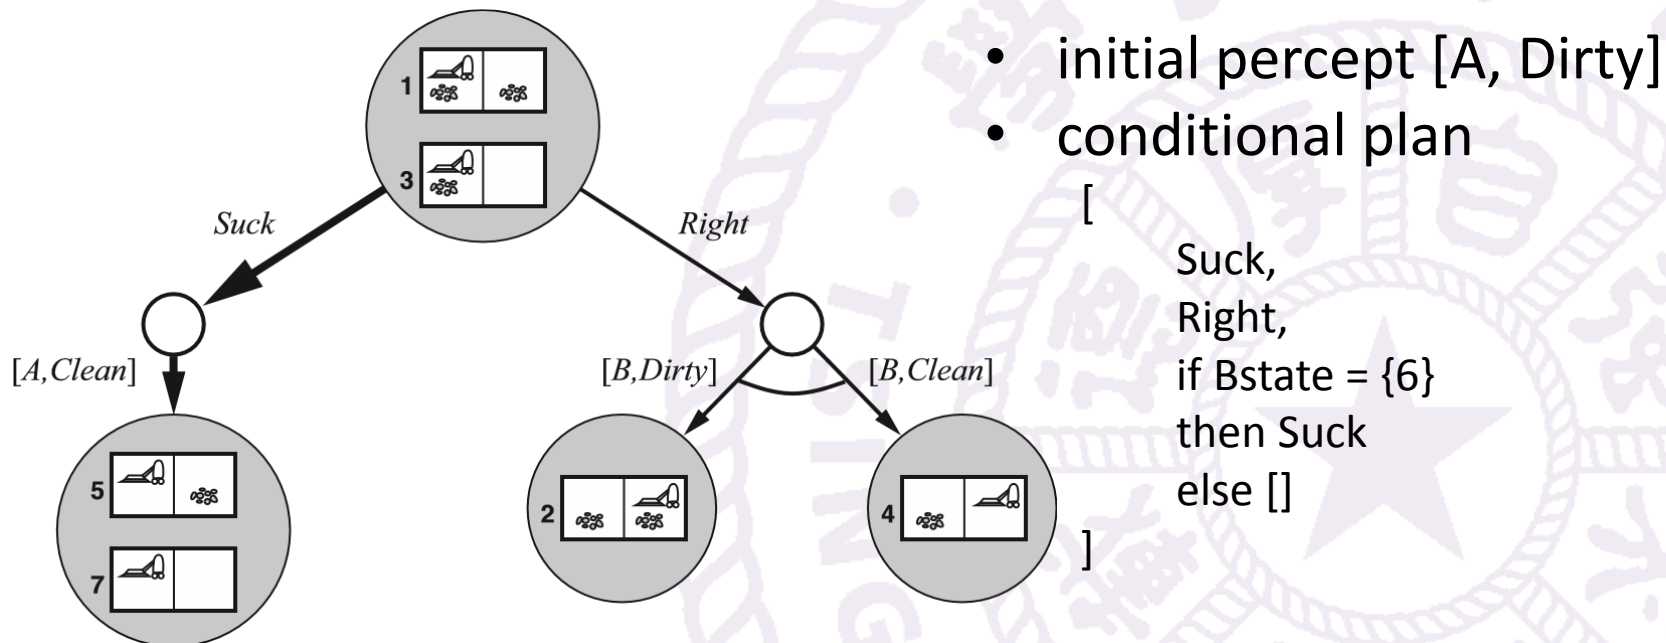


- b. Slippery world

local-sensing vacuum worlds

Solving Partially Observable Problems

- A nondeterministic belief-state problem with
 - *ACTIONS, RESULTS, GOAL-TEST, STEP-COSTS*
- AND-OR search algorithm can be applied to derive a solution



Summary

- Local Search
 - Hill climbing
 - Simulated annealing
 - Genetic algorithms
 - Continuous space
- Search with Uncertainty
 - Nondeterministic Actions
 - Partial Observations

谢谢！

